

Bhavya Vashisht

+91 9810589918 | Delhi NCR | bhavyavashisht119@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Education

Thapar Institute Of Engineering and Technology	Patiala
Bachelor of Engineering in Electronics and Communication Engineering	2022 – 2026
CGPA: 8.75	
Shalom Hills International School	Gurugram
CLASS XII (PCM with CS) CBSE 95.8%	2020 – 2021

Experience

Yum! Brands – Digital Food Platform	Gurugram
Machine Learning Engineer Intern	January, 2026 – Present
– Built production recurring agent scheduling on Kubernetes CronJobs — REST schedule API with 7 recurrence patterns, CronJob lifecycle manager (<code>concurrencyPolicy: Forbid</code>), and a worker that orchestrates agent/team invocation and writes run records to Snowflake. – Designed Redis-backed one-time scheduling with at-least-once delivery, an idempotency guard across horizontally scaled worker replicas, and a single-replica Snowflake→Redis sync coordinator so DB read load stays constant as workers scale — periodic reconcile for reliable production execution after Redis failover or wipe. – Shipped async LLM cost tracking: non-blocking emit after each chat turn, Redis ZSET job queue, background worker with retry/backoff, and Snowflake model-pricing cache — chat responses never wait on cost calculation. – Built production K8s deployment infrastructure for agents and multi-agent teams — unified REST deploy API (create/save/redeploy/undeploy), Helm lifecycle with Snowflake status tracking, init-container GitLab preload for runtime actions, and deployment-scoped ingress + asset URLs for isolated pod routing. – Integrated Model Context Protocol (stdio + Streamable HTTP); fixed a production race condition in shared MCP config via request-scoped deep-copy across 3 execution paths. – Optimized a MongoDB conversation-list API via metadata-only endpoint + compound index — payload 1.6 MB → 7 KB (99.5%), latency 6–15s → sub-500ms.	

Skills

Languages: Python, Go, C, SQL
Frameworks & Libraries: FastAPI, PydanticAI, MCP, JWT
Databases & Caching: MongoDB, PostgreSQL, Redis
Tools: Git, Docker, Kubernetes, Helm, GitLab CI/CD, Linux, SQLC, Goose
Coursework: Data Structures & Algorithms, OS, DBMS, OOP, AI, Computer Architecture
Certifications: HTTP Clients in Go, Learn Linux, Learn OOP in Python, Learn Go

Projects

Chirpy Go, PostgreSQL, JWT, SQLC, Goose	GitHub
– Built a REST API in Go for a micro-blogging platform with JWT auth, bcrypt hashing, and middleware for logging, validation, and error recovery. – Designed a PostgreSQL schema with SQLC + Goose (5 migrations, 20+ queries); added Chirpy Red premium upgrades via webhooks.	
Gator Go, PostgreSQL, SQLC, Goose, RSS	GitHub
– Built a CLI RSS aggregator in Go with multi-user feed follow/unfollow, PostgreSQL persistence via SQLC + Goose, and a continuous aggregation worker that round-robins feed fetches on a configurable ticker. – Designed idempotent post ingestion via unique constraints so re-fetching never duplicates entries; type-safe queries generated from SQL by sqlc.	
JobOps Pipeline Python, Google Sheets API, asyncio/httpx, LLM APIs	GitHub
– Built a daily job-ops pipeline (~8.6k LOC): <code>morning.py</code> orchestrates Sheet sync → parallel fetch across 15+ job boards and ATS APIs → fresher/geo/salary filters → tiered scoring → bidirectional Google Sheets CRM. – Implemented dedup by URL plus company/role/location collapse; optional LLM JD enrichment (Gemini/OpenAI) extracts stack/YOE and auto-picks resume variant; cached in <code>.jd_cache.json</code>. – Added outreach step: Hunter.io contact lookup, per-role cold-mail drafts, and company referral DMs — review-before-send workflow (never auto-sends).	